

Baby Name Popularity

Francisco Rivera
frivera@college.harvard.edu

Mark Chamberlain
mark_chamberlain@college.harvard.edu

November 15, 2018

Abstract

The U.S. Census publishes the number of babies born to each first name each year at the country-level. It has previously (Hahn and Bentley, 2003) been suggested these trends could be explained by a simple imitation model where parents copy baby names from the previous year at random. However, we find these trends exhibit dramatic up- and down-swings unexplainable by this simple model. We propose a segmented-population model to explain the observed data. Our model succeeds in fitting the data and is very suggestive of an agent-based generative process, but suffers from numerical instabilities in parameter choices and

1 The Data

Our data comes from the U.S. Census.¹ The dataset allows us to observe the number of births in the country each year by first name and gender with the exception of names for which there were fewer than five births.

The data requires little processing, but because our aim is to model changes in popularity rather than population growth, we normalize each data entry as a percentage of the total number of births that year. To provide a sense of what some typical trends look like, we plot the popularity of five random popular² names in Figure 1.

On the basis of inspecting many realizations of Figure 1, for which the depicted trends are representative, we can take away some stylized facts,

- The popularity of a name in many cases follows an exponential-like growth from obscurity, a peak, and then a decay back to obscurity.
- There appears to be a practical upper-bound to how high the peak gets, but many (unpopular) names never get close to that upper-bound.
- As “Deborah” shows, the down-swing need not be as fast as the up-swing.

¹<https://catalog.data.gov/dataset/baby-names-from-social-security-card-applications-national-level-data>

²Defined as exceeding 0.5% of all births on any year.




Figure 1: Popularity evolution of five random names

2 An Imitation Model

We follow the gist of Hahn and Bentley (2003). At its core, the model is straightforward: we assume that every time period (i.e. each year) every baby is born by choosing a random baby from the previous time period and copying that name. Hahn and Bentley (2003) highlights the main strengths of this model which are two-fold,

- It is extremely parsimonious. There is some nuance surrounding initialization, but they find that when initializing to individually different names in epoch 1, the model converges after only a couple hundred time periods. Beyond that, the model has no parameters to fit and its assumption of consecutive time-period copying is sensible.
- The cross-sectional distribution of baby-name popularity matches the power-law predicted by this imitation model. That is, in any given year, the distribution of the popularity of different names is consistent with the dynamics of this model.

We accept these strengths at face value, but question whether this model is sufficient to model the inter-temporal aspect as well. To test this, we introduce some notation. For the i^{th} name let $p_t^{(i)}$ represent the proportion of the number of births in year t to name i . Our key model assumption is that,

$$p_{t+1}^{(i)} \sim \frac{1}{N_{t+1}} \cdot \text{Binom}(N_{t+1}, p_t^{(i)}).$$

This formula also explains why we are not concerned with the initial conditions of the model: conditional on the previous time period's distribution, the model dictates a probability distribution over the following period.

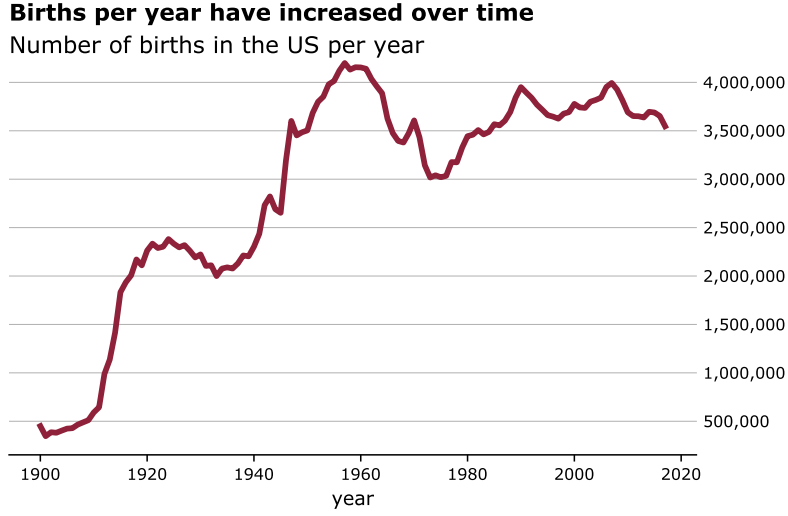


Figure 2: Births per year

In particular, the only parameter in this equation is the number of births which is easily observable. We plot the number of births per years in Figure 2. This graphs shows us that while births have generally been increasing over the course of the 20th century, the number of births has remained approximately the same order of magnitude, and for the latter part of the century can be reasonably approximated as $N_t \approx 3.5 \times 10^6$.

This allows us to calculate

$$\text{Var}(p_{t+1}^{(i)} | p_t^{(i)}) = \frac{p_t^{(i)}(1 - p_t^{(i)})}{N_{t+1}}$$

In particular, if p is moderately large to say that the distribution of $p_{t+1}^{(i)} | p_t^{(i)}$ is approximately distributed as,

$$\mathcal{N}\left(p_t^{(i)}, \frac{p_t^{(i)}(1 - p_t^{(i)})}{N_{t+1}}\right)$$

For small values of p_t , $1 - p_t^{(i)} \approx 1$ and our standard deviation will be approximately,

$$\frac{\sqrt{p_t^{(i)}}}{\sqrt{N_{t+1}}} \approx \frac{\sqrt{p_t^{(i)}}}{10^3}.$$

For $p_t^{(i)} = 1\%$ this means that the year-on-year standard deviation becomes 10^{-4} . If $p_t^{(i)} = 0.01\% = 10^{-4}$ the year-on-year standard deviation becomes 10^{-5} . And if $p_t = 10^{-6}$ the year-on-year standard deviation will be approximately 10^{-6} . This illustrates an important consequence of this model: the because the standard deviation scales as approximately the square root of the popularity of the names (for small enough popularities, where the $1 - p_t^{(i)}$ is approximately

1), we have that the standard deviation as a multiple of the proportion is bigger for smaller probabilities.

Put more concretely, we are unsurprised if a very unpopular name doubles in popularity from one year to another, but are quite surprised if this happens for a more popular name. However, does this prediction survive the data?

Figure 3 illustrates a resounding no. In it, we depict 100 randomly generated paths from the imitation model superimposed on a starting condition for the popularity of the name “Amanda,” as well as the actual realized trend. We see that according to this model, the boom in popularity of the name “Amanda” from 0.5% of the popularity to 2.5% of the yearly baby population would be near-impossible. This is not a cherry-picked example: these swings in popularity constantly appear across the trends of different names.




Figure 3: Imitation model predictions

Figure 3 runs the model using the precise yearly births and generating binomial random variables. What we get looks an awfully lot like Brownian motion, though, which we can explain by our previous observation. We had previously noted that the binomial tends toward a Normal by the Central Limit Theorem, and when

$$\sqrt{\text{Var}(p_{t+1}^{(i)} | p_t^{(i)})} \ll p_t$$

we have that,

$$p_{t+1} \approx p_t \implies \text{Var}(p_{t+1}^{(i)} | p_t^{(i)}) \approx \text{Var}(p_{t+2}^{(i)} | p_{t+1}^{(i)})$$

giving us an approximation to the model prediction for uncertainty k years into the future as,

$$\text{Var}(p_{t+k}^{(i)} | p_t^{(i)}) \approx k \cdot \frac{p_t^{(i)}(1 - p_t^{(i)})}{N_{t+1}}.$$

However, we face the fact that our trends look distinctly unlike Brownian motion, and we must consider a different model to explain the shape of these trends.

3 Segmenting the population

Another way to view the shortcomings of the imitation model or a model like it is that the evolution of name popularity appears to behave differently on the up-swing than the down-swing. This became starkly inconsistent with the imitation model in the amount of variance we see from year-to-year in name popularity, but even if the variance matched, we consistently see a boom followed by a bust in name popularity, which remains unexplainable by this model.

This suggests that there is something different in the adoption of a name at different points in the stylized life-cycle of popularity. One way to model this is by segmenting the population, akin to the work on fads done in Bergman et al. (2012). We posit that there are two sub-populations: an “in” group and an “out”-group. “In”-group membership is desirable, and members of this group try to distinguish themselves through names distinctive of the group, yet at the same time not over-used.

Members of the out-group also wish to adopt desirable in-group names. However, we posit that the out-group will only be able to react to lagged information of name popularity.

Mathematically, then, in our most general form, we are suggesting if $p_{\text{in}}^{(i)}(t)$ (respectively $p_{\text{out}}^{(i)}(t)$) are the proportion of the in-group (respectively out-group) at time t given name i , then our evolution equations are,

$$\begin{aligned}\frac{dp_{\text{in}}^{(i)}}{dt} &= \alpha_{\text{in}} \cdot p_{\text{in}}^{(i)} \left(K_{\text{in}} - p_{\text{in}}^{(i)} - \beta_{\text{in}} \cdot p_{\text{out}}^{(i)} \right) \\ \frac{dp_{\text{out}}^{(i)}(t)}{dt} &= \alpha_{\text{out}} \cdot p_{\text{out}}^{(i)} \left(p_{\text{in}}^{(i)}(t - \tau) - \beta_{\text{out}} \cdot p_{\text{out}}^{(i)}(t - \tau) \right)\end{aligned}$$

In this equations, the α parameters represent a sensitivity of the system to its pressures, i.e. it calibrates how quickly members of the in and out-group react to incentives to shift toward or away a name. K_{in} represents the maximal sustainable size of the proportion of the in-group with the name in the absence of any out-group adopters. The β parameters capture aversion to names being used by members of the out-group, and τ represents the lag with which members of the out-group can observe name trends.

Qualitatively, we expect the life-cycle of a name to have five stages:

1. A name becomes representative of the in-group. This perturbation need not be large, and could be result of random noise.
2. As members of the in-group see this name as representative of the in-group, they begin to adopt it more and more. Because the out-group only observed lagged popularity, the in-group popularity can boom unencumbered by out-group adoption.
3. The out-group realizes the name is representative of the in-group and starts adopting it as well. This is the hey-day of the name popularity, when both the in-group and the out-group adopt it.
4. The in-group responds to the out-group’s adoption and begins using it less. The out-group continues using it because it doesn’t realize the in-group has moved away from it.

5. The out-group finally realizes that the name is no longer being used by the in-group and moves away from it as well. Now the name fades back into obscurity.

3.1 Numerical simulation

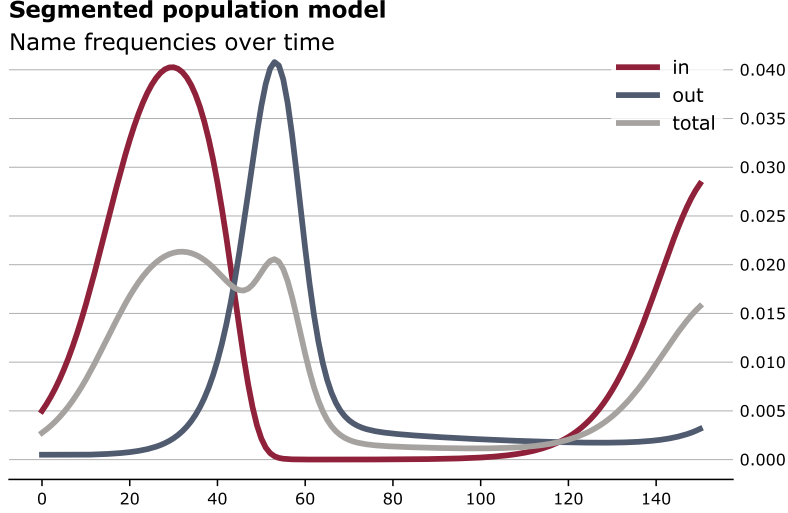


Figure 4: Segmented model simulation

In order to better understand the behavior these equations impose, we can simulate the evolution of the in-group, the out-group, and the total popularity of a name over time for the evolution of the model, for instance we depict one such configuration in Figure 4. In it, we have assumed that the in-group and the out-group are equally represented (each constitutes 50% of the total population), but this need not be the case. If we make the out-group be the majority of the population, we simply get that the popularity of the name in the entire population very closely follows the popularity of the name in the out-group.

Given our ability to numerically simulate, we can then understand the effect of different parameters by varying them and seeing how they affect the curve. For instance, consider the parameter α_{out} which measures the sensitivity to incentives in the out-group. The bigger it is, the faster out-group members move toward desirable names and away from undesirable names. We would thus expect bigger values of α_{out} to correspond to faster boom-bust life-cycles, as well as to higher absolute levels of popularity (because the lag is held constant, a fast response to incentives means a bigger increase in the name popularity before it starts going down. Both of these predictions hold in the data, which we can see in Figure 5.

We expect the lag in our model to have similar effects on the maximum popularity of a name: the longer the lag, the more popular the name will become because the out-group continues adopting the name for a longer time before realizing it is no longer desirable. However, we expect an opposite effect on the

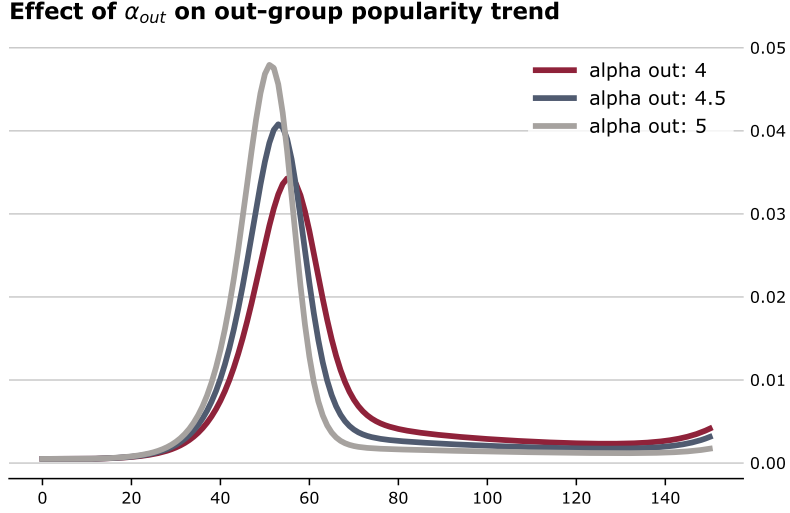


Figure 5: Effect of α_{out} on popularity trend

boom-bust lifecycle duration: a longer lag will lead to a longer duration of the adoption period. These conclusions are confirmed in Figure 6.

Thus, when we fit the model to an observed name’s popularity trend we are able to match the empirical trend with a theoretical one by fitting these parameters to the shape of the curve. Most notably, the attributes of the curve that are matched include,

- The maximum height of popularity.
- The duration of the up-swing
- The duration of the down-swing
- Whether the curve has two peaks

3.2 Fitting the model

Thus far we have explored the effect of different parameters on the model output, gaining an intuition for how the model behaves. However, if we wish to use this model for prediction we need to be able to fit to data. To do this, we use an optimization library (`scipy`) to minimize our squared error loss using the Nelder-Mead optimization procedure on all parameters except the lag, which because it must be an integer is trained through a grid-search instead.

We find numerical instability for random initialization of our parameter vector. This is a complication to batch training, and we discuss it more in the evaluation of our model. However, when we initialize to hand-tuned parameters, the optimization routine converges to locally optimal parameters, matching the data reasonably well. For instance, we display the empirical trend for the popularity of “Amanda” superimposed with the model fit in Figure 7.

Parameters vary significantly across fitting the trends of different names, which is consistent with there being significant heterogeneity in trends across

Effect of lag on out-group popularity trend




Figure 6: Effect of lag τ on popularity trend

names. For “Amanda” our optimal fit tells us that the proportion of the in-group is approximately 25%, the α s are between 1 and 2, the lag comes out as 3 years, and the β come out as approximately 6.

4 Evaluating the model

The model has a number of strengths. At the forefront of these strengths is the interpretability of the model underpinnings. The agents in the model (members of the in-group and out-group) behave in an intuitively appealing way. That is, the model sheds light on what we believe the underlying behavior is being driven by.

In addition, the model has the flexibility to match the data we observe. Unlike the simple imitation model which under-predicts the variance, this model can closely track the wide variety of trends that we observe in the data. The parameters of the model are also interpretable, since we can reason about their effects intuitively as well as through numerical simulation.

Finally, the model makes predictions about the data that are confirmed. For instance, the model suggests that trends should have one or two peaks depending on the proportion of the in-group, and this is empirically confirmed.

However, the model also has some significant weaknesses. Foremost among these is the number of parameters that need to be estimated. The large number of parameters make us wary of overfitting. Even though the parameters are interpretable, they are not directly observable, which also makes the fitting process complicated. From our work, we see that the fitting process is sensitive to initial conditions, which complicates the use of our model for a generic trend.

At a higher level, there are also important limitations to what our model captures. Perhaps most relevantly, the model has no underlying theory to why different names should have different trends / parameters. A more complete

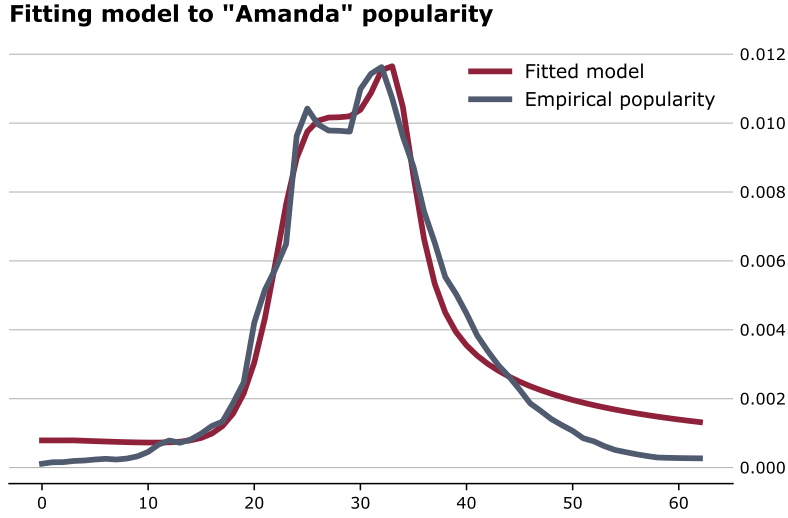


Figure 7: Fit of model to “Amanda” trend

model would have to answer the different question of why different names behave differently.

5 Future work

One shortcoming previously noted was the large number of parameters in the model, which decrease parsimony, lead to over-fitting concerns, and complicate the fitting process. One way to improve the model would be to capture the same qualitative behavior with fewer distinct parameters. One way we attempted to do this was by setting the α, β parameters for the in and out-groups to be the same. Unfortunately, that constraint led to oscillatory behavior counterfactual to observed trends. Nevertheless, we believe it should be possible with a smart constraint to simplify the parameter space while maintaining the model’s ability to capture the data.

A strength we tout is the interpretability of model’s the underlying generation process. However, future work is needed to determine whether this generation process is actually indicative of reality. In particular, our model hypothesizes the existence of in- and out-groups with in-group booms preceding out-group booms in a name’s popularity. It would be reasonable to hypothesize that members of the in-group are more famous / rich / “elite” than out-group members. Thus, a testable prediction of the model is that trends among these “elites” should lead country-wide popularity trends. Finding data on naming trends among these demographics would be a pivotal test for our model.

Finally, we return to an initial ambition of the project, which was to also model the spatial component of name popularity. In this work, we have modeled the temporal aspect by looking at name popularity over time. It would be interesting to take the generative model we have posited and adapt it to also make predictions about how name popularity will diffuse geographically.

References

- Bergman, M. et al. (2012). When a fad ends: An agent-based model of imitative behavior. *Computing in Economics and Finance*. Available: <http://www.uh.edu/margo/paper.pdf>. Accessed.
- Hahn, M. W. and Bentley, R. A. (2003). Drift as a mechanism for cultural change: an example from baby names. *Proceedings of the Royal Society of London B: Biological Sciences*, 270(Suppl 1):S120–S123.